

Physics-Informed Map-Conditioned UE Localization

Coarse-to-fine Transformer posterior with differentiable radio-map physics

Nir Tzur

January 19, 2026

Story arc

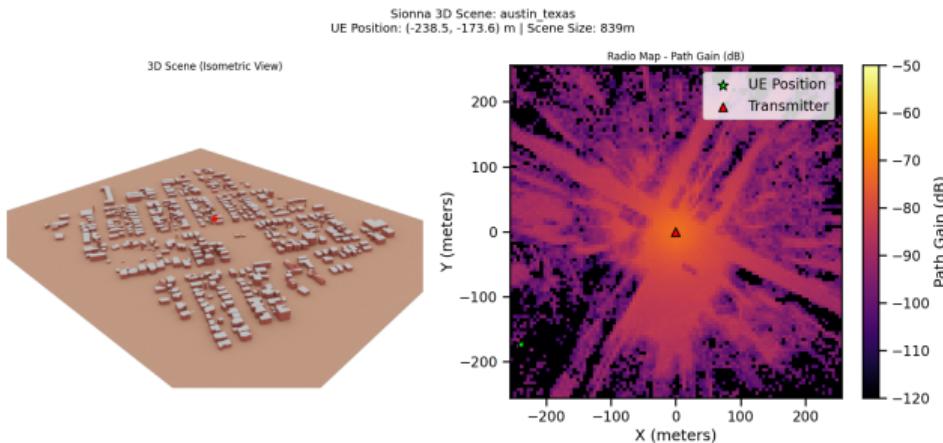
- Problem framing: localization is multi-modal and physics-constrained.
- Pipeline: OSM to scenes to radio maps to training and inference.
- Model: dual encoder, cross-attention fusion, coarse-to-fine posterior.
- Evidence: qualitative results, error statistics, and training dynamics.

Problem and goals

- Input: sparse, irregular radio measurements across multiple protocol layers.
- Challenge: inverse mapping is ill-posed and multi-modal in urban scenes.
- Goal: map-conditioned probabilistic posterior with calibrated uncertainty.
- Constraint: predictions must remain consistent with propagation physics.

End-to-end pipeline

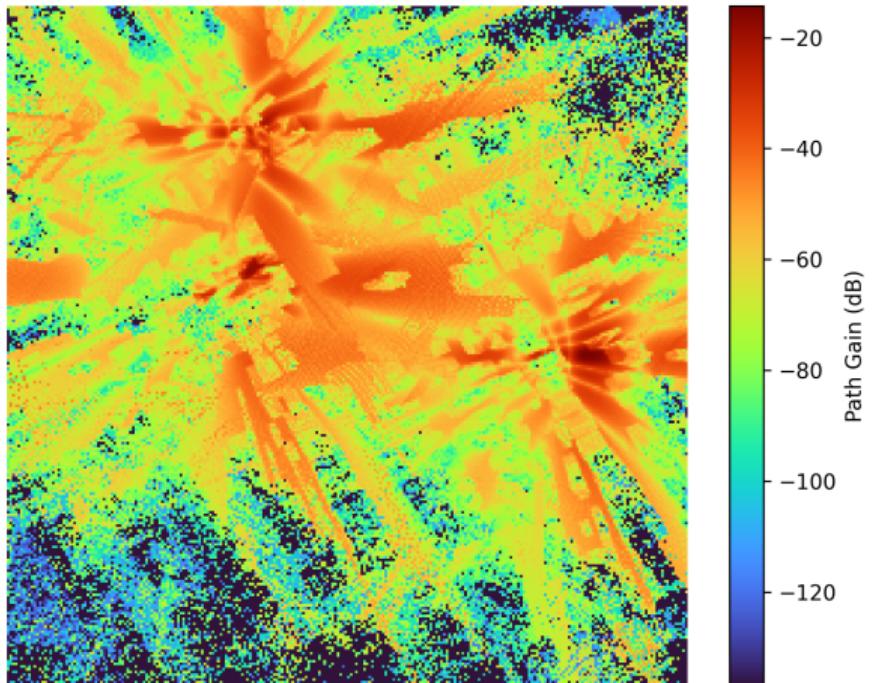
- OSM and GIS data define scene geometry and semantics.
- Differentiable ray tracing yields radio maps and measurements.
- Train a transformer to amortize inference over UE position.
- Optional refinement improves low-confidence predictions.



Scene and radio map from Sionna RT.

Map context for localization

Radio Map (Signal Propagation)



Semantic Map (Buildings)

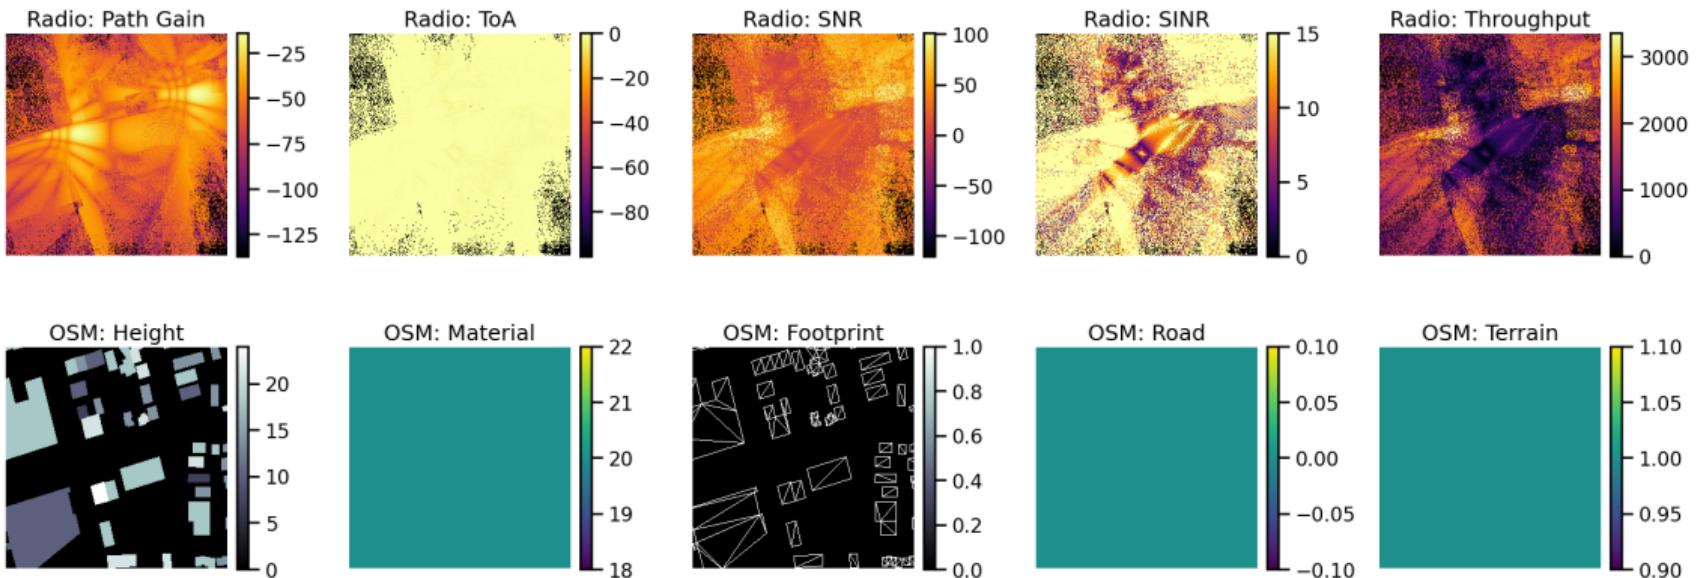


Radio propagation map and semantic building context.

Inputs and map channels

- Radio features: RT, PHY/FAPI, and MAC/RRC measurements.
- Map features: radio channels (path gain, ToA, SNR, SINR, throughput) and OSM channels (height, material, footprint, road, terrain).

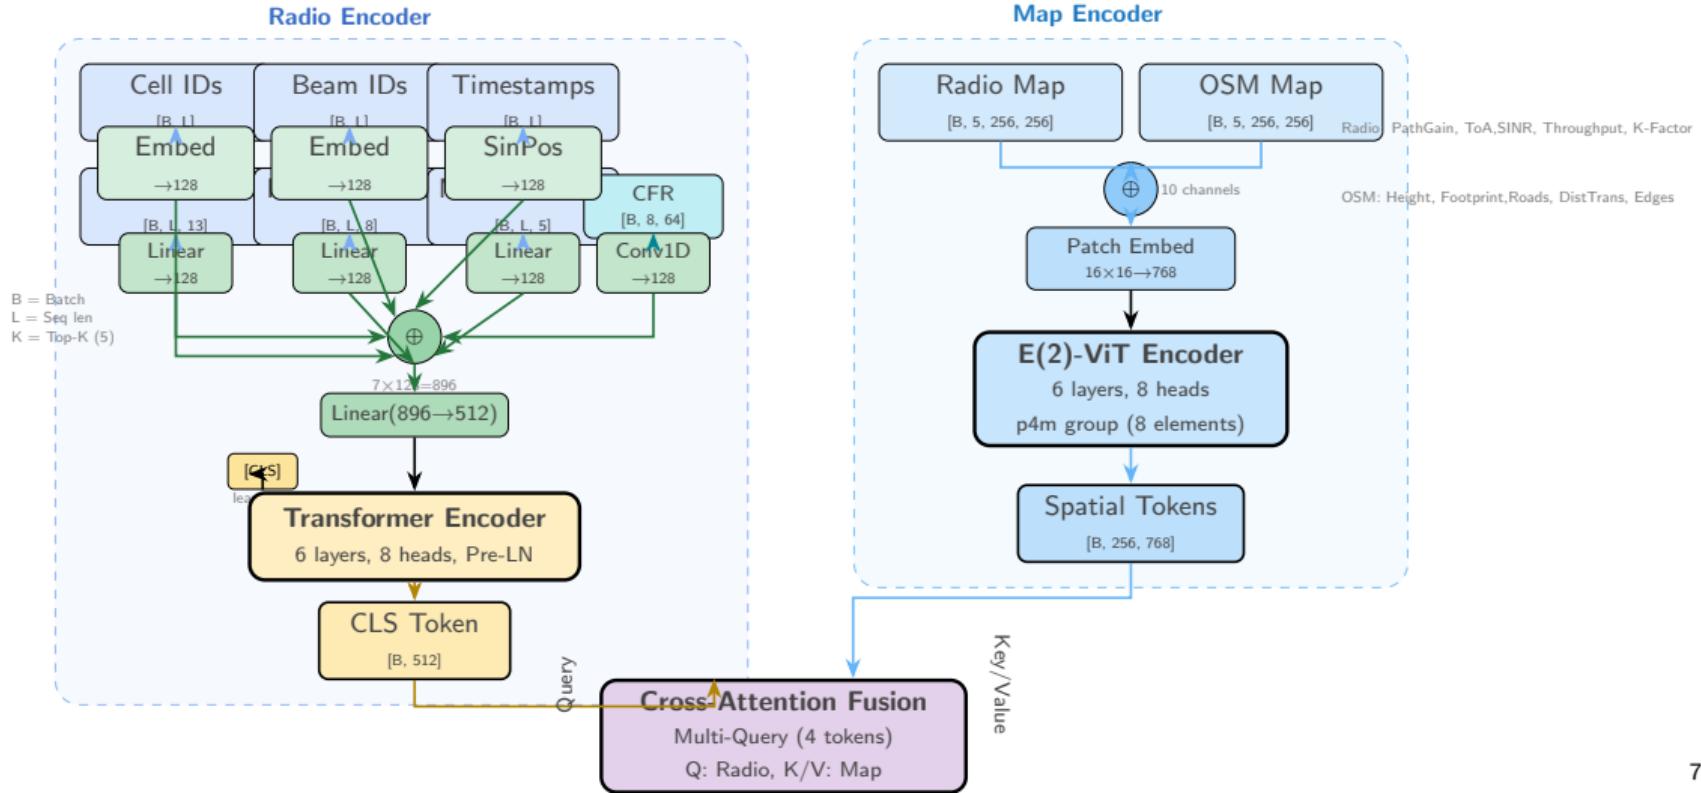
Map Layers (Sample 0)



Model architecture

UE Localization Model Architecture

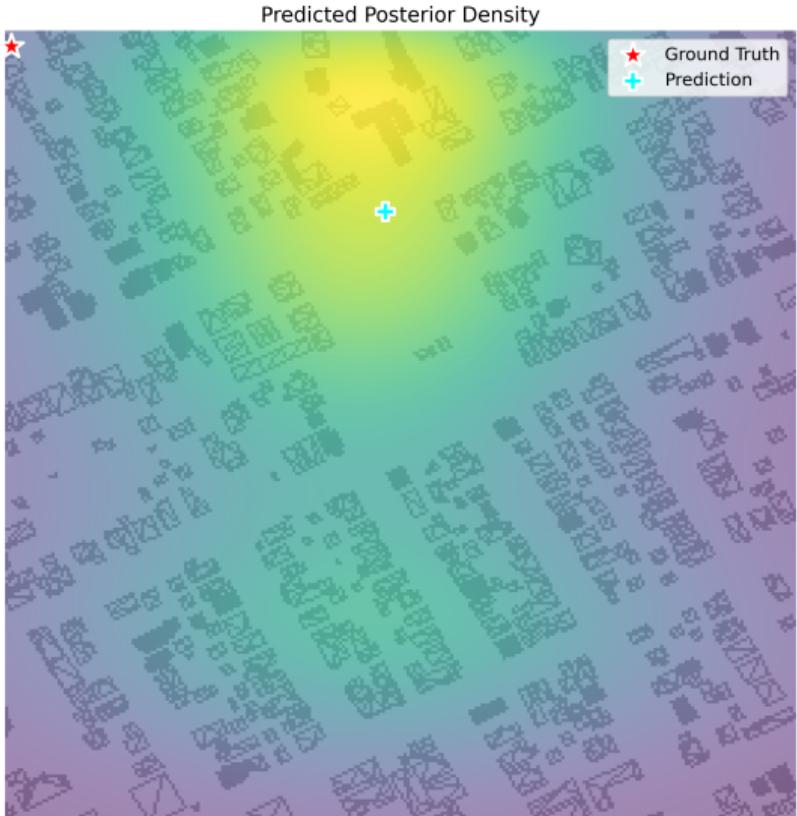
Coarse-to-Fine Map-Conditioned Transformer



Coarse-to-fine posterior

- Coarse grid predicts a heatmap over spatial cells.
- Top-K cells refine to Gaussian components.
- Mixture posterior encodes multi-modality and uncertainty.

$$p(\mathbf{y}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{y}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$



Training objective and physics loss

Total objective

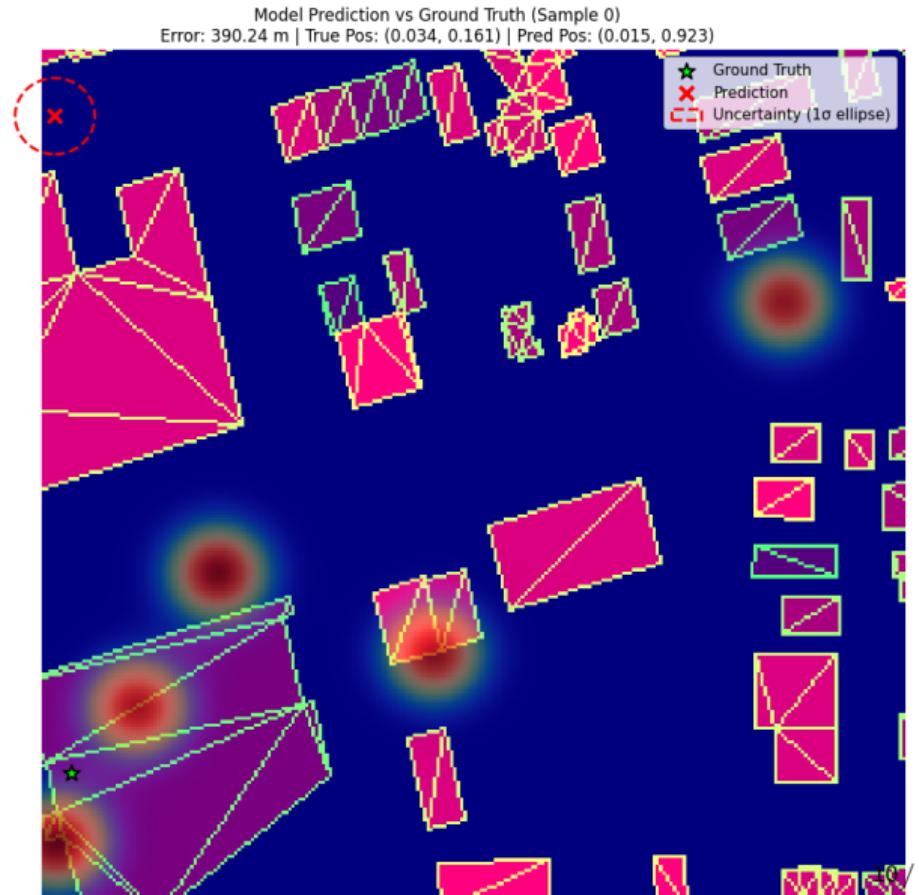
$$\mathcal{L} = \mathcal{L}_{\text{coarse}} + \lambda_f \mathcal{L}_{\text{fine}} + \lambda_p \mathcal{L}_{\text{phys}}$$

- Coarse: cross-entropy over grid cells.
- Fine: Gaussian NLL with heteroscedastic uncertainty.
- Physics: differentiable lookup of radio maps via bilinear sampling.

$$\mathcal{L}_{\text{phys}} = \sum_f w_f \|m_f^{\text{obs}} - R_f(\hat{\mathbf{y}})\|^2$$

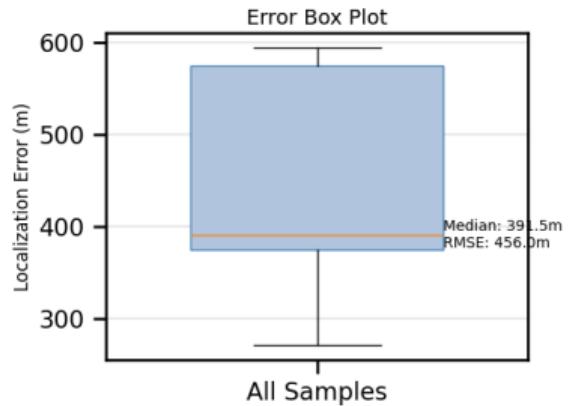
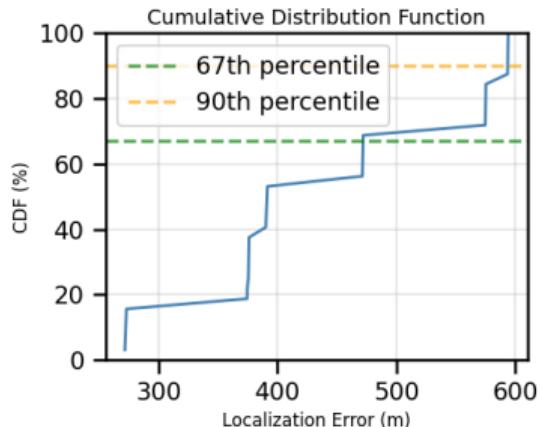
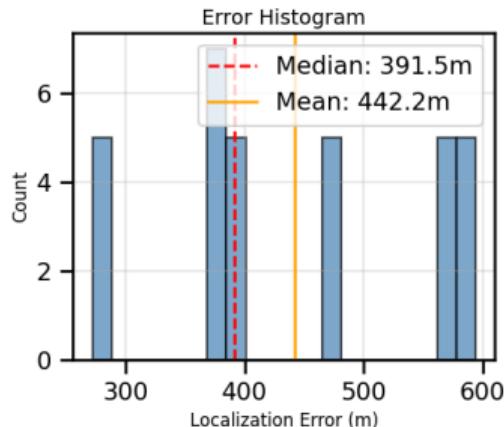
Inference and refinement

- Predict heatmap and top-K candidate cells.
- Refine to continuous positions with uncertainty.
- Optional MAP refinement combines network posterior with radio-map likelihood.



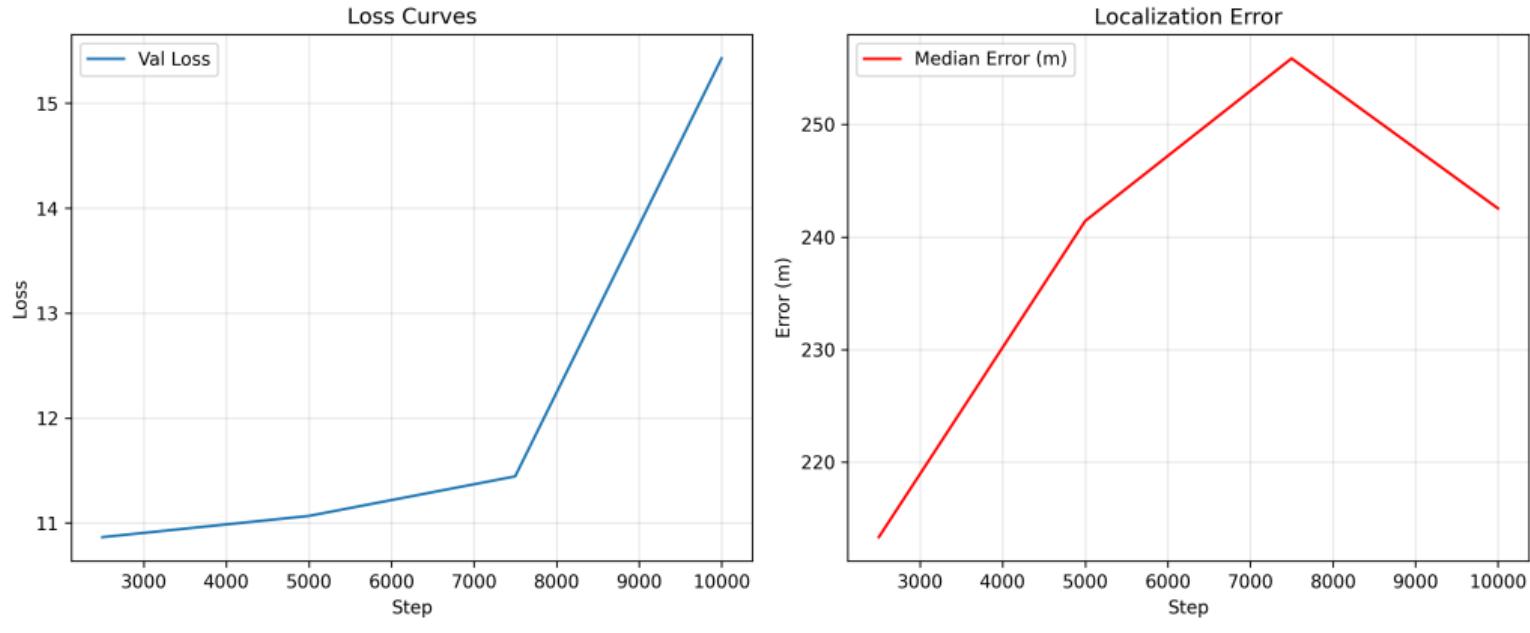
Evaluation metrics

Batch Error Analysis



Error histogram, CDF, and box plot for batch evaluation.

Training dynamics



curves and median localization error over steps.

Project status and integration

- Scene generation from OSM.
- Differentiable ray tracing and dataset generation.
- Dual-encoder transformer with fusion.
- Training and evaluation pipeline.
- Streamlit web demo.

Milestone	Status
Scene generation	Complete
Data generation	Complete
Model	Complete
Training	Complete
Web UI	Complete

Conclusion and next steps

- Map-conditioned transformer yields a calibrated, multi-modal posterior.
- Physics-informed loss aligns predictions with radio propagation.
- Pipeline is end-to-end: OSM to scenes to training to inference.
- Next: larger-scale city runs, ablations, and uncertainty calibration.